

The screenshot displays a Jupyter Notebook titled "sleep_detector.ipynb" with the following code cells:

```
[1] import os
import cv2
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
[2] dataset_path = "/content/drive/MyDrive/dataset"
```

```
[3] def load_data(dataset_path):
    images = []
    labels = []
    classes = os.listdir(dataset_path)
    for i, class_name in enumerate(classes):
        class_path = os.path.join(dataset_path, class_name)
        for image_name in os.listdir(class_path):
            image_path = os.path.join(class_path, image_name)
            image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
            image = cv2.resize(image, (32, 32))
            images.append(image)
            labels.append(i)
    images = np.array(images, dtype=np.float32) / 255.0
    labels = np.array(labels)
    return images, labels
```

```
[4] images, labels = load_data(dataset_path)
```

```
[5] # 데이터 증강
datagen = ImageDataGenerator(
    rotation_range=1,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True
)

augmented_images = []
augmented_labels = []

for i in range(images.shape[0]):
    image = images[i]
    image = np.expand_dims(image, axis=-1)
    image = np.expand_dims(image, axis=0)
    label = labels[i]
    for _ in range(10): # 각 이미지에만 10장의 추가 생성
        x_augmented, y_augmented = datagen.flow(image, [label], batch_size=1)
        augmented_images.append(np.squeeze(x_augmented))
        augmented_labels.append(y_augmented[0])
    break

augmented_images = np.array(augmented_images)
augmented_labels = np.array(augmented_labels)
```

```
[6] # 데이터 분할
x_train, x_test, y_train, y_test = train_test_split(augmented_images, augmented_labels, test_size=0.2, random_state=42)
```

```
[7] model = models.Sequential([
    layers.Conv2D(128, (3, 3), activation='relu', input_shape=(32, 32, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(256, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dense(3, activation='softmax') # 클래스 수만큼의 출력
])
```

```
[8] model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 128)	1280
max_pooling2d (MaxPooling2D)	(None, 15, 15, 128)	0
conv2d_1 (Conv2D)	(None, 13, 13, 128)	147584
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 128)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	295168
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 256)	1048832
dense_1 (Dense)	(None, 3)	771

Total params: 1493635 (5.70 MB)
Trainable params: 1493635 (5.70 MB)
Non-trainable params: 0 (0.00 Byte)

```
[9] # 모델 컴파일
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
[10] # 모델 학습
history = model.fit(x_train, y_train, epochs=200, batch_size=32, validation_data=(x_test, y_test))

print(history.history)
```

```
import matplotlib.pyplot as plt
```

Epoch 161/200

```
2/2 [=====] - 0s 146ms/step - loss: 2.0853e-04 - accuracy: 1.0000 - val_loss: 2.4497 - val_accuracy: 0.5833
```

Epoch 162/200

```
2/2 [=====] - 0s 143ms/step - loss: 2.0724e-04 - accuracy: 1.0000 - val_loss: 2.4536 - val_accuracy: 0.5833
```

Epoch 163/200

```
2/2 [=====] - 0s 159ms/step - loss: 2.0419e-04 - accuracy: 1.0000 - val_loss: 2.4439 - val_accuracy: 0.5833
```

Epoch 164/200

```
2/2 [=====] - 0s 150ms/step - loss: 2.0002e-04 - accuracy: 1.0000 - val_loss: 2.4475 - val_accuracy: 0.5833
```

Epoch 165/200

```
2/2 [=====] - 0s 167ms/step - loss: 1.9739e-04 - accuracy: 1.0000 - val_loss: 2.4397 - val_accuracy: 0.5833
```

Epoch 166/200

```
2/2 [=====] - 0s 156ms/step - loss: 1.9440e-04 - accuracy: 1.0000 - val_loss: 2.4362 - val_accuracy: 0.5833
```

Epoch 167/200

```
2/2 [=====] - 0s 154ms/step - loss: 1.9229e-04 - accuracy: 1.0000 - val_loss: 2.4346 - val_accuracy: 0.5833
```

```

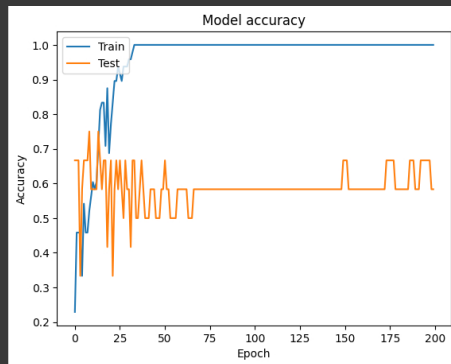
Epoch 166/200
2/2 [=====] - 0s 165ms/step - loss: 1.9032e-04 - accuracy: 1.0000 - val_loss: 2.4368 - val_accuracy: 0.5833
Epoch 169/200
2/2 [=====] - 0s 145ms/step - loss: 1.8818e-04 - accuracy: 1.0000 - val_loss: 2.4426 - val_accuracy: 0.5833
Epoch 170/200
2/2 [=====] - 0s 147ms/step - loss: 1.8618e-04 - accuracy: 1.0000 - val_loss: 2.4555 - val_accuracy: 0.5833
Epoch 171/200
2/2 [=====] - 0s 162ms/step - loss: 1.8314e-04 - accuracy: 1.0000 - val_loss: 2.4535 - val_accuracy: 0.5833
Epoch 172/200
2/2 [=====] - 0s 186ms/step - loss: 1.8045e-04 - accuracy: 1.0000 - val_loss: 2.4650 - val_accuracy: 0.5833
Epoch 173/200
2/2 [=====] - 0s 140ms/step - loss: 1.7777e-04 - accuracy: 1.0000 - val_loss: 2.4618 - val_accuracy: 0.5833
Epoch 174/200
2/2 [=====] - 0s 156ms/step - loss: 1.7577e-04 - accuracy: 1.0000 - val_loss: 2.4611 - val_accuracy: 0.6667
Epoch 175/200
2/2 [=====] - 0s 153ms/step - loss: 1.7376e-04 - accuracy: 1.0000 - val_loss: 2.4621 - val_accuracy: 0.6667
Epoch 176/200
2/2 [=====] - 0s 152ms/step - loss: 1.7186e-04 - accuracy: 1.0000 - val_loss: 2.4620 - val_accuracy: 0.6667
Epoch 177/200
2/2 [=====] - 0s 165ms/step - loss: 1.6973e-04 - accuracy: 1.0000 - val_loss: 2.4653 - val_accuracy: 0.6667
Epoch 178/200
2/2 [=====] - 0s 147ms/step - loss: 1.6768e-04 - accuracy: 1.0000 - val_loss: 2.4676 - val_accuracy: 0.6667
Epoch 179/200
2/2 [=====] - 0s 140ms/step - loss: 1.6560e-04 - accuracy: 1.0000 - val_loss: 2.4726 - val_accuracy: 0.5833
Epoch 180/200
2/2 [=====] - 0s 188ms/step - loss: 1.6341e-04 - accuracy: 1.0000 - val_loss: 2.4767 - val_accuracy: 0.5833
Epoch 181/200
2/2 [=====] - 0s 152ms/step - loss: 1.6132e-04 - accuracy: 1.0000 - val_loss: 2.4809 - val_accuracy: 0.5833
Epoch 182/200
2/2 [=====] - 0s 139ms/step - loss: 1.5983e-04 - accuracy: 1.0000 - val_loss: 2.4843 - val_accuracy: 0.5833
Epoch 183/200
2/2 [=====] - 0s 154ms/step - loss: 1.5767e-04 - accuracy: 1.0000 - val_loss: 2.4828 - val_accuracy: 0.5833
Epoch 184/200
2/2 [=====] - 0s 200ms/step - loss: 1.5699e-04 - accuracy: 1.0000 - val_loss: 2.4799 - val_accuracy: 0.5833
Epoch 185/200
2/2 [=====] - 1s 260ms/step - loss: 1.5381e-04 - accuracy: 1.0000 - val_loss: 2.4868 - val_accuracy: 0.5833
Epoch 186/200
2/2 [=====] - 1s 233ms/step - loss: 1.5234e-04 - accuracy: 1.0000 - val_loss: 2.4968 - val_accuracy: 0.5833
Epoch 187/200
2/2 [=====] - 1s 237ms/step - loss: 1.5066e-04 - accuracy: 1.0000 - val_loss: 2.4991 - val_accuracy: 0.6667
Epoch 188/200
2/2 [=====] - 1s 224ms/step - loss: 1.4919e-04 - accuracy: 1.0000 - val_loss: 2.5010 - val_accuracy: 0.6667
Epoch 189/200

```

```

[11] plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



```

[12] # 모델 저장
model.save("image_classification_model.h5")

```

```

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead `saving_api.save_model(

```